

PSTAT 126


Regression Analysis

Laura Baracaldo & Rodrigo Targino


Lecture 1
Introduction to Linear Models

Introduction

Regression Analysis: Statistical method that allows the exploration of the relationship between two or more variables.



Linear Regression Models



Introduction

- Y : Response, outcome, output, dependent variable.
- $(X_1, \dots, X_p)^T$: vector of predictors, inputs, independent or explanatory variables.

Goals:


- 1 Prediction of future or unmeasured responses given values of the predictors.
- 2 Assessment of the effect/relationship between explanatory variables and response. Existence of causality?

Cars example

[Link to help](#)

X : Speed Y : Distance


```
par(mar=c(3,4,0.5,1))  
plot(cars, pch=18)
```



Cars example

X : Speed Y : Distance

```
par(mar=c(3,4,0.5,1))  
plot(cars, pch=18)  
abline(lm( dist ~ speed, data=cars), col="blue")
```




Cheddar example

```
## Warning: package 'faraway' was built under R version 4.2.2
```

Link to help

X : Hydrogen sulfide concentration Y : Taste score

```
library(faraway)
par(mar=c(3,4,0.5,1))
plot(cheddar$H2S, cheddar$taste, pch=18)
```




Cheddar example

X : Hydrogen sulfide concentration

Y : Taste score

```
library(faraway)
par(mar=c(3,4,0.5,1))
plot(cheddar$H2S, cheddar$taste, pch=18)
abline(lm( taste ~ H2S, data=cheddar), col="blue")
```




Pressure example

[Link to help](#)

X : Temperature Y : Pressure


```
par(mar=c(3,4,0.5,1))  
plot(pressure, pch=18)
```



Pressure example


X : Temperature Y : Pressure

```
par(mar=c(3,4,0.5,1))  
plot(pressure, pch=18)  
abline(lm( pressure ~ temperature, data=pressure), col="red")
```




Regression vs Classification

- Regression: Y is continuous. Numerical values (e.g. price, blood pressure, confirmed COVID-19 cases).
- Classification Y discrete labels. Categorical values. (e.g. survived/died, digit 0-9, if Bitcoin price is going up tomorrow).



General Workflow




1. Data: We collect data for all n subjects in the study, and identify Y and X :

$$Y = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix}$$

$$X = \begin{bmatrix} x_{11} & x_{12} & \dots & x_{1p} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n1} & x_{n2} & \dots & x_{np} \end{bmatrix}$$

General Workflow




1. Data: We collect data for all n subjects in the study, and identify Y and X :

$$Y = \begin{bmatrix} y_1 \\ y_2 \end{bmatrix}$$

$$X = \begin{bmatrix} x_{11} & x_{12} & \dots & x_{1p} \\ x_{21} & x_{22} & \dots & x_{2p} \end{bmatrix}$$

General Workflow



1. Data: We collect data for all n subjects in the study, and identify Y and X :

$$Y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$$X = \begin{bmatrix} x_{11} & x_{12} & \dots & x_{1p} \\ x_{21} & x_{22} & \dots & x_{2p} \\ \vdots & \vdots & \dots & \vdots \\ x_{n1} & x_{n2} & \dots & x_{np} \end{bmatrix}$$

- $Y \in \mathbb{R}^n, X \in \mathbb{R}^{n \times p}$

Regression Model

$$Y = f(X_1, X_2, \dots, X_p) + \epsilon$$

- ϵ : Error/noise
- f is unknown and must be learned from observations (Y, X_1, \dots, X_p)
- Questions about f . Continuous? Smooth?
- f must be restricted. In this course: We assume f is linear.

Linear Regression Model

We assume the response is written as linear representation of predictors:

$$Y = \beta_0 + \beta_1 X_1 + \dots \beta_p X_p + \epsilon$$

- **Estimation:** The goal is to estimate unknown parameters $\beta_0, \beta_1, \dots, \beta_p$.

Linear Regression Model

- **Linear:**

$$Y = \beta_0 + \beta_1 \log(X_1) + \dots \beta_p X_p^2 + \epsilon = \beta_0 + \beta_1 W_1 + \dots \beta_p W_p + \epsilon$$

- **NOT Linear:**

$$Y = \beta_0 + X_1^{\beta_1} + \dots \beta_p X_p + \epsilon$$

Matrix Representation

$$\begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} 1 & x_{11} & \dots & x_{1p} \\ 1 & x_{21} & \dots & x_{2p} \\ \vdots & \vdots & \dots & \vdots \\ 1 & x_{n1} & \dots & x_{np} \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{bmatrix}$$

$$Y = X\beta + \epsilon$$

Model Assessment

- ➊ Prediction: accurately predict future response given predictors.
- ➋ Inference: assess the quality of our predictions and (or) estimation.
- ➌ Diagnostics: what if the linear model assumptions go wrong? how can we tell?
- ➍ Model selection: find the “best” linear model for response given predictors.

Next...

- Set up your R working environment.
- Lab sessions and TAs' office hours start next week.
- Office hours:
 - Prof. Baracaldo: TBD
 - Prof. Targino: Fridays 1.30pm - 3.30pm OG1230 or on Zoom
<https://fgv-br.zoom.us/j/95030256507>